

# DDP & AllReduce

Distributed Data Parallel

2025

# Распределенный блок

- GPU clusters
- **DDP & AllReduce** <----- **you are here**
- FSDP
- TP & CP & EP & PP ...

# Сегодняшняя лекция

- Distributed Training
- Data Parallel
- AllReduce
- Overlapping with optim step
- Unsynced GD

# Сегодняшний семинар

- nccl tests
- ddp training
- optimizer
- memory + perf profiles



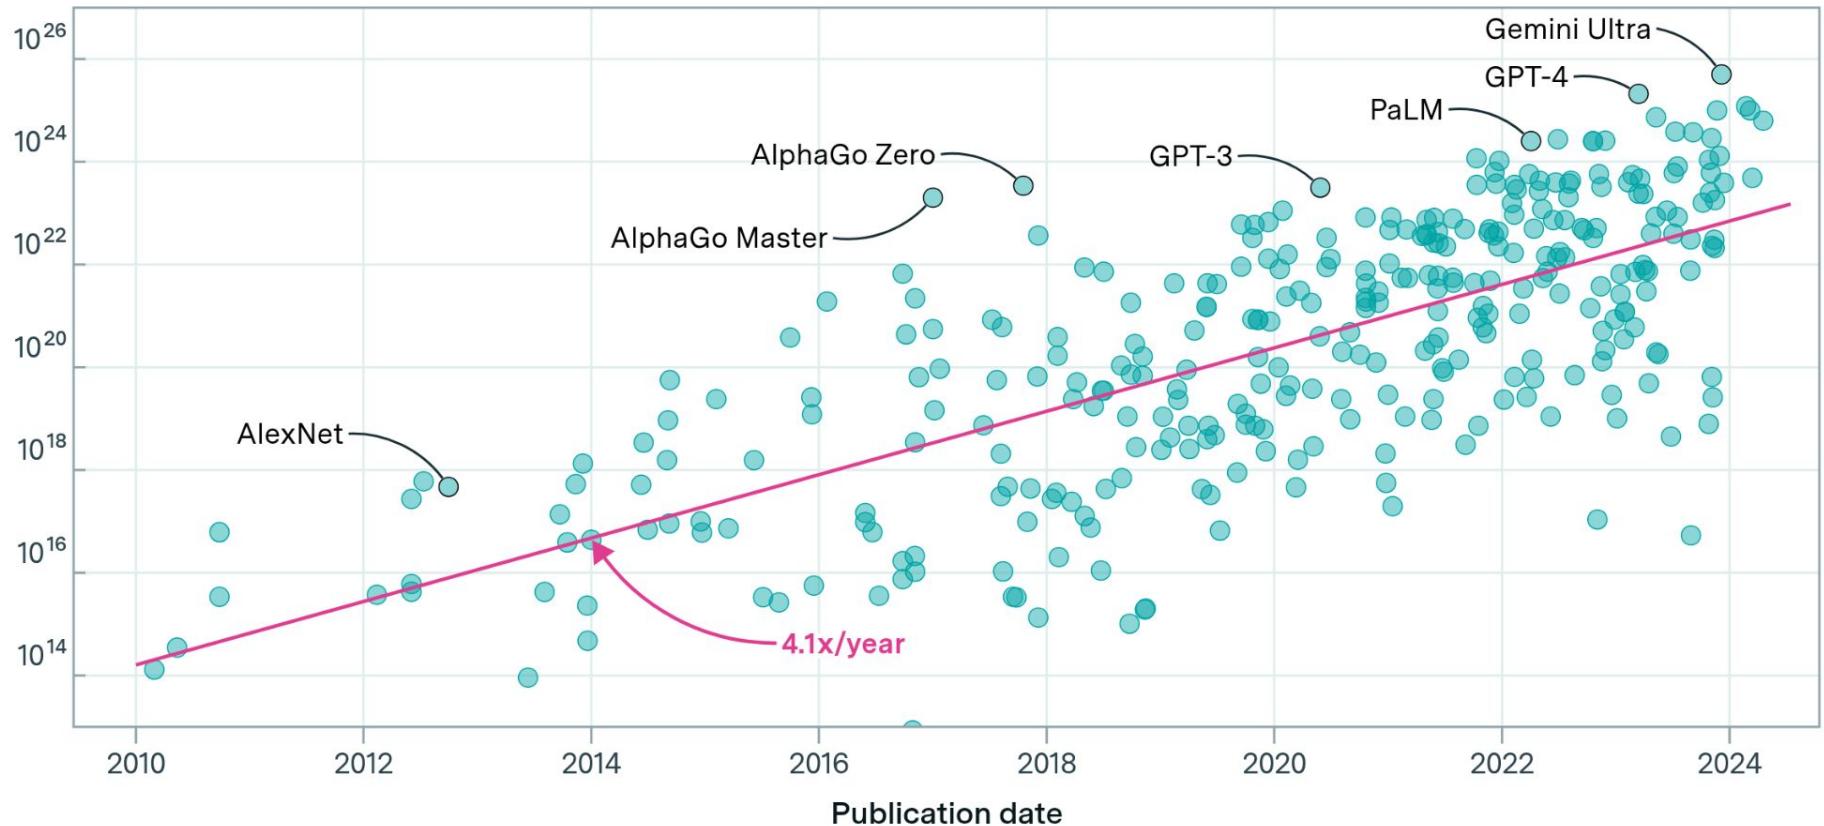
# Distributed Training

# Training compute of notable models



Training compute (FLOP)

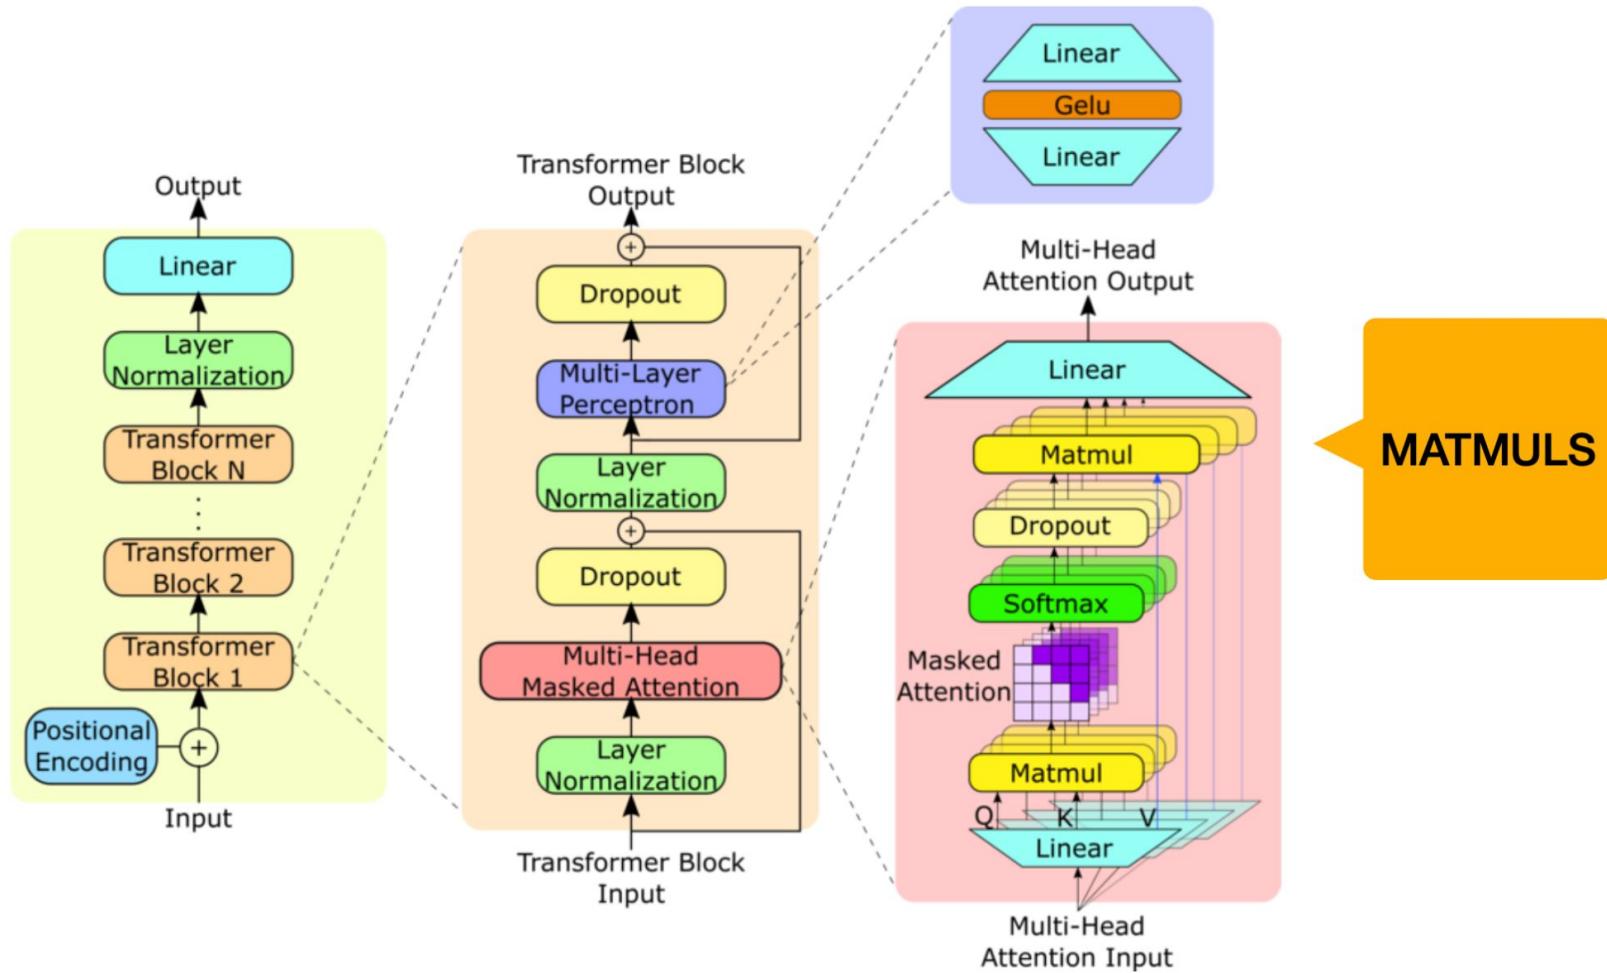
333 models



CC-BY

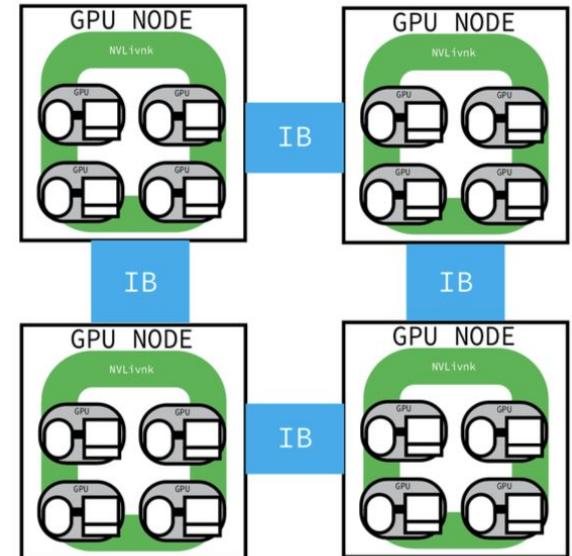
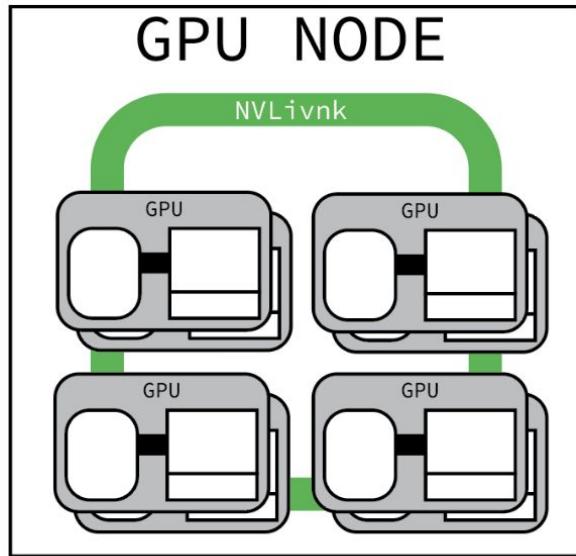
epoch.ai

# Recap: transformer



# Recap: GPU clusters

- GPU объединяются в сервера а.к.а “ноды” с быстрой сетью
- Ноды объединяются в кластера



# Parallelisms

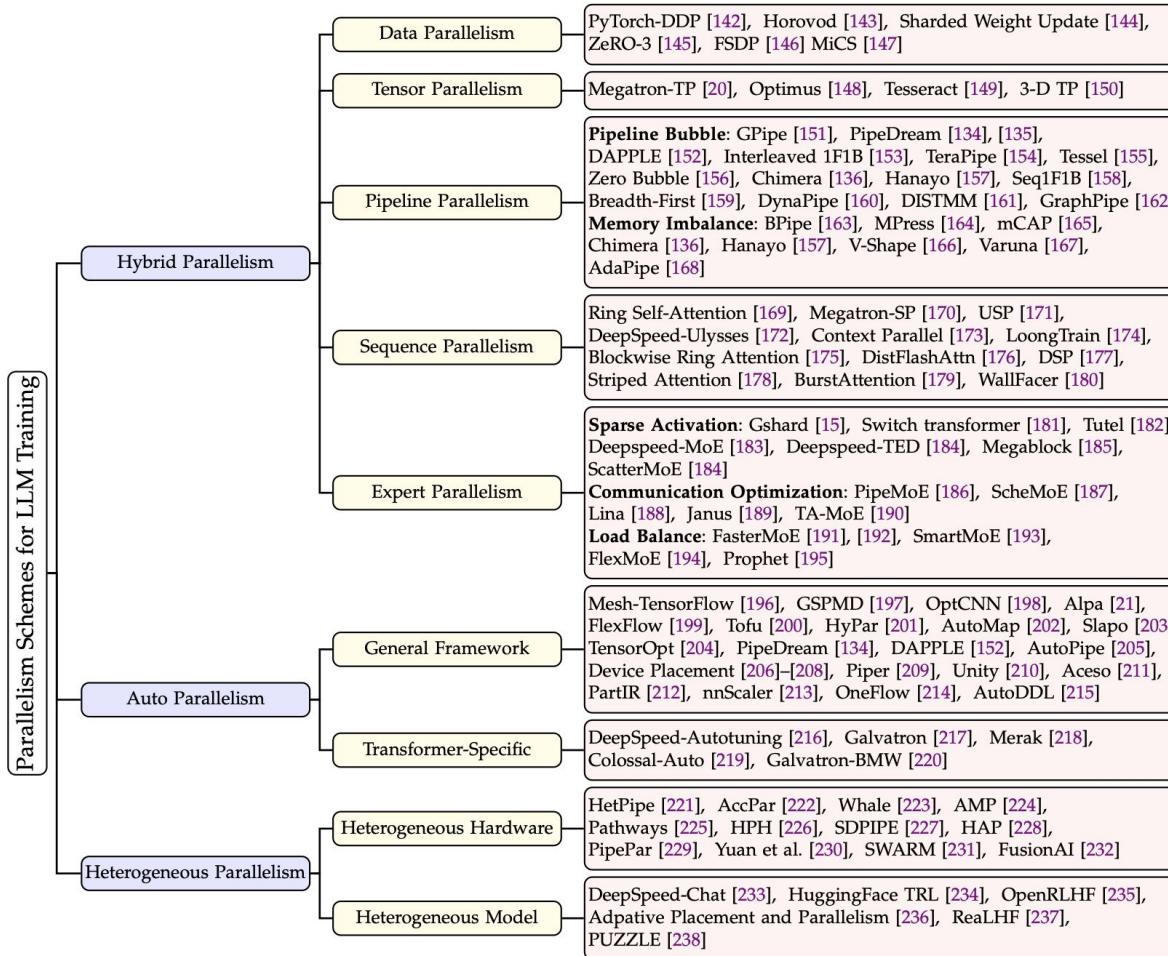
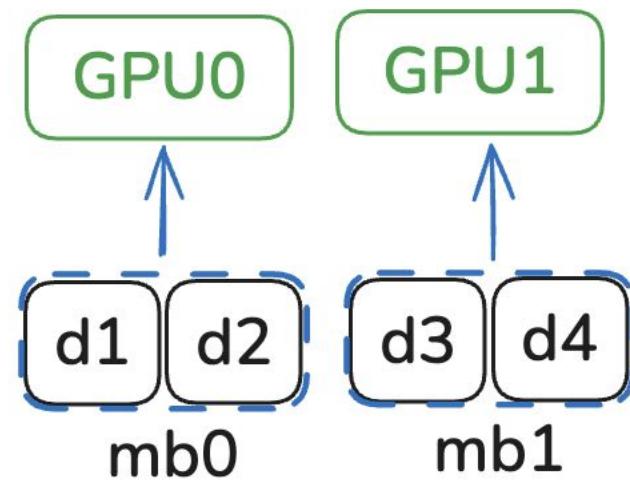
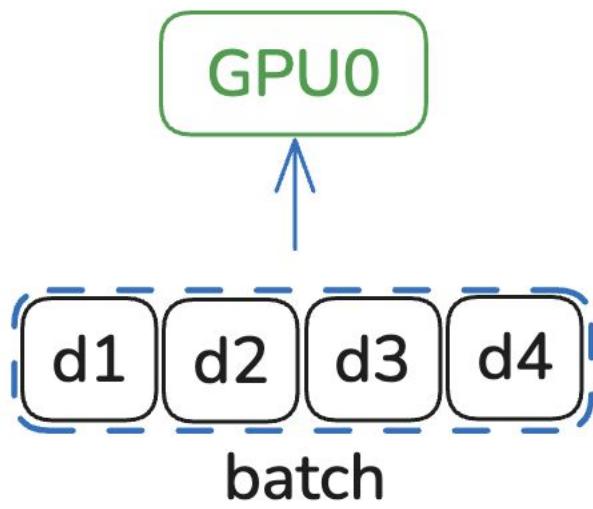


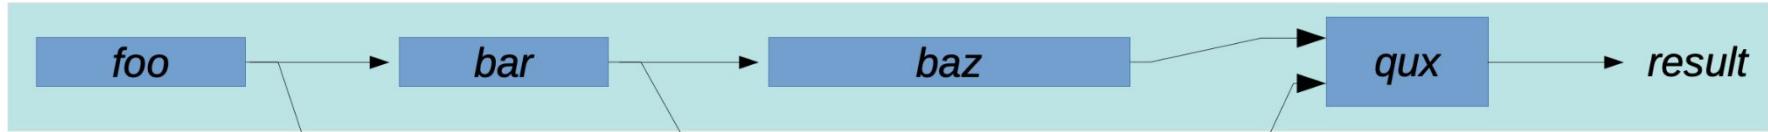
Fig. 7: Studies on parallelism schemes for distributed LLM training.



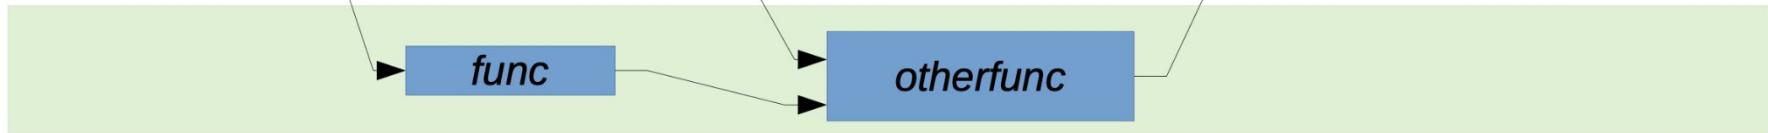
# Data Parallel

# Channel / Pipe

**Process A:**



**Process B:**



**Channel (pipe):**

- Communication in  $O(\text{message size})$
- Asynchronous read/write

\*Слайды EffDL



A meme stolen from somewhere

# AllReduce NCCL

# AllReduce

```
torch.distributed.all_reduce(tensor, op=<RedOpType.SUM: 0>, group=None,  
async_op=False)
```

[\[source\]](#)

Reduces the tensor data across all machines in a way that all get the final result.

After the call `tensor` is going to be bitwise identical in all processes.

Complex tensors are supported.

## Parameters:

- **tensor** ([Tensor](#)) – Input and output of the collective. The function operates in-place.
- **op** (*optional*) – One of the values from [torch.distributed.ReduceOp](#) enum. Specifies an operation used for element-wise reductions.
- **group** ([ProcessGroup](#), *optional*) – The process group to work on. If None, the default process group will be used.
- **async\_op** ([bool](#), *optional*) – Whether this op should be an async op

## Returns:

Async work handle, if `async_op` is set to True. `None`, if not `async_op` or if not part of the group

# AllReduce

## Examples

---

```
>>> # All tensors below are of torch.int64 type.  
>>> # We have 2 process groups, 2 ranks.  
>>> device = torch.device(f"cuda:{rank}")  
>>> tensor = torch.arange(2, dtype=torch.int64, device=device) + 1 + 2 * rank  
>>> tensor  
tensor([1, 2], device='cuda:0') # Rank 0  
tensor([3, 4], device='cuda:1') # Rank 1  
>>> dist.all_reduce(tensor, op=ReduceOp.SUM)  
>>> tensor  
tensor([4, 6], device='cuda:0') # Rank 0  
tensor([4, 6], device='cuda:1') # Rank 1
```

```
>>> # All tensors below are of torch.cfloat type.  
>>> # We have 2 process groups, 2 ranks.  
>>> tensor = torch.tensor(  
...     [1 + 1j, 2 + 2j], dtype=torch.cfloat, device=device  
... ) + 2 * rank * (1 + 1j)  
>>> tensor  
tensor([1.+1.j, 2.+2.j], device='cuda:0') # Rank 0  
tensor([3.+3.j, 4.+4.j], device='cuda:1') # Rank 1  
>>> dist.all_reduce(tensor, op=ReduceOp.SUM)  
>>> tensor  
tensor([4.+4.j, 6.+6.j], device='cuda:0') # Rank 0  
tensor([4.+4.j, 6.+6.j], device='cuda:1') # Rank 1
```

# AllReduce

TABLE V  
STEPS IN ONE LOOP ITERATION OF NCCL RING ALLREDUCE

Step Index	NCCL Primitive
0	send
1 to $k - 2$	recvReduceSend
$k - 1$	recvReduceCopySend
$k$ to $2k - 3$	recvCopySend
$2k - 2$	recv

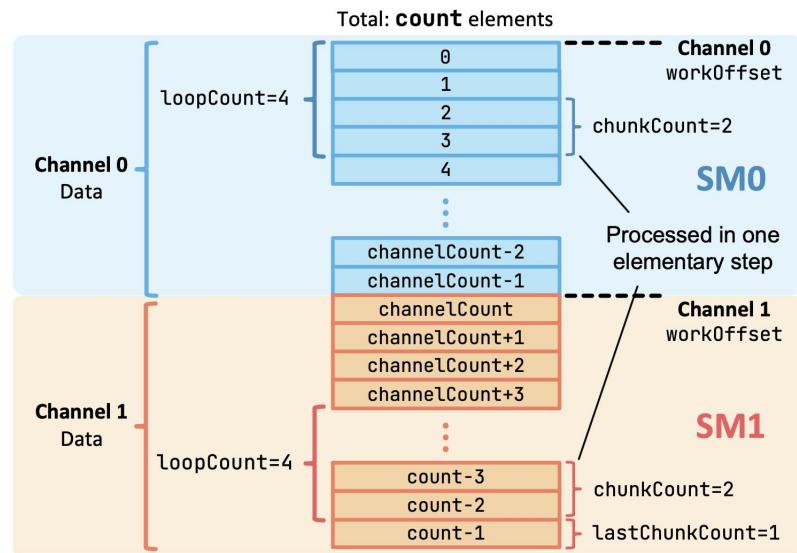
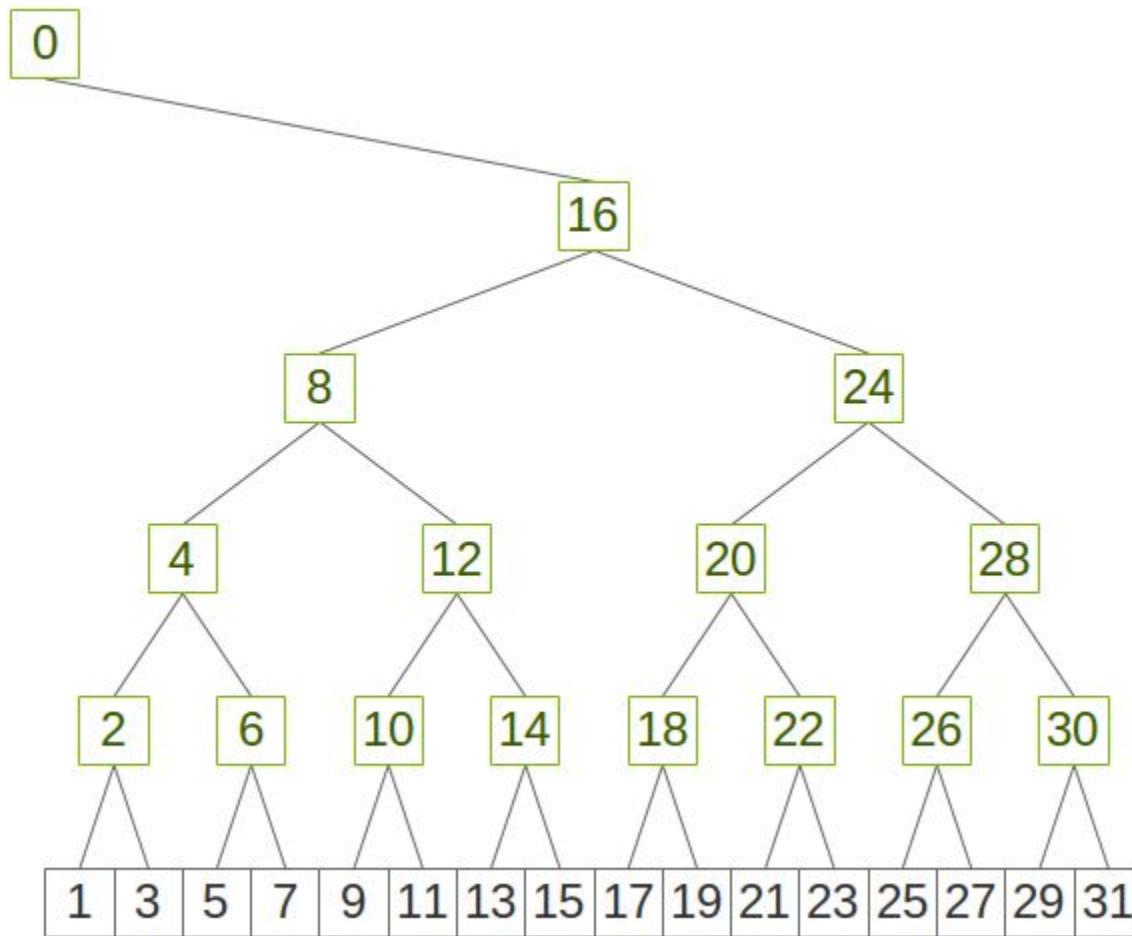
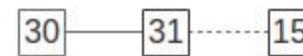
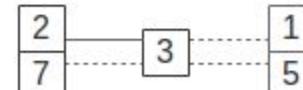
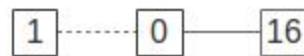
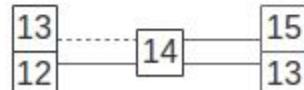
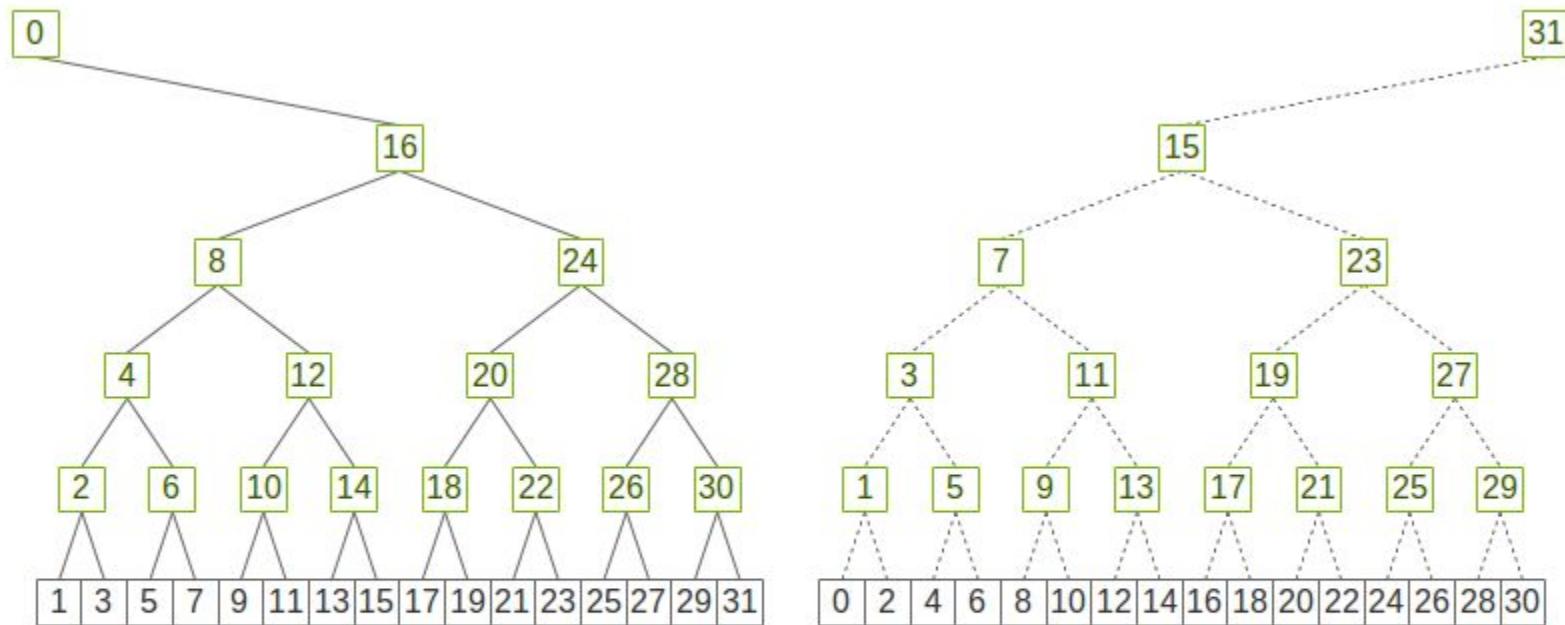


Fig. 3. Visualization of NCCL's data partitioning strategy across communication channels and loop iterations.

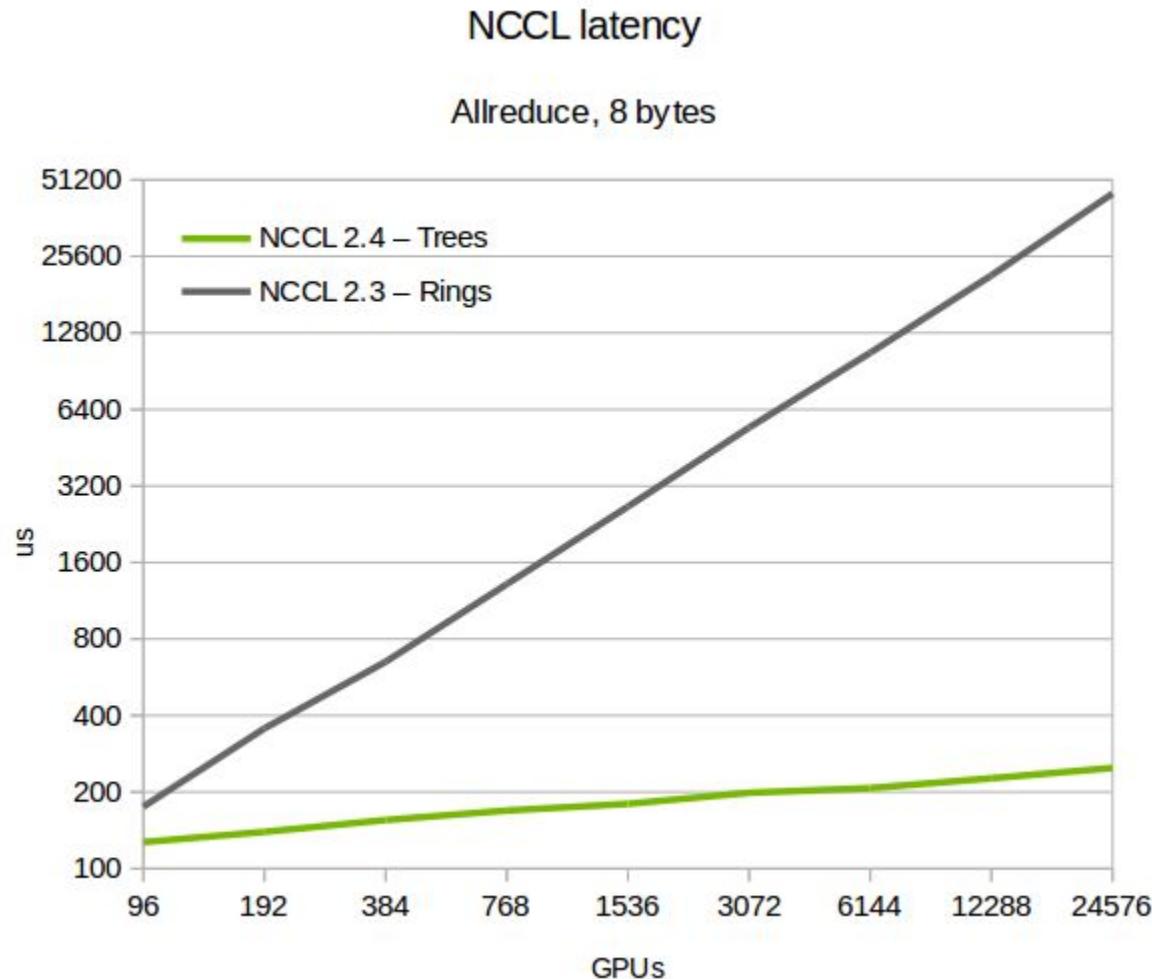
# Tree AllReduce



# Tree AllReduce



# Ring vs Tree AllReduce



# NCCL tests

nccl-tests Public Watch 28

master 3 Branches 39 Tags Go to file Add file Code

stephenmsachs and AddyLaddy Check if sufficient GPUs are available · abc4677 · 2 weeks ago · 105 Commits

doc	doc: add all2all factor	last year
src	Check if sufficient GPUs are available	2 weeks ago
verifiable	Make verifiable a DSO and add NAME_SUFFIX support	6 months ago
.gitignore	Fix #43 : Add .gitignore for build dir	5 years ago
LICENSE.txt	Initial commit	8 years ago
Makefile	makefile: remove extra space	2 years ago
README.md	Modified warmup to run for more message sizes	2 months ago

README BSD-3-Clause license

## NCCL Tests

These tests check both the performance and the correctness of [NCCL](#) operations.

# NCCL benchmarking

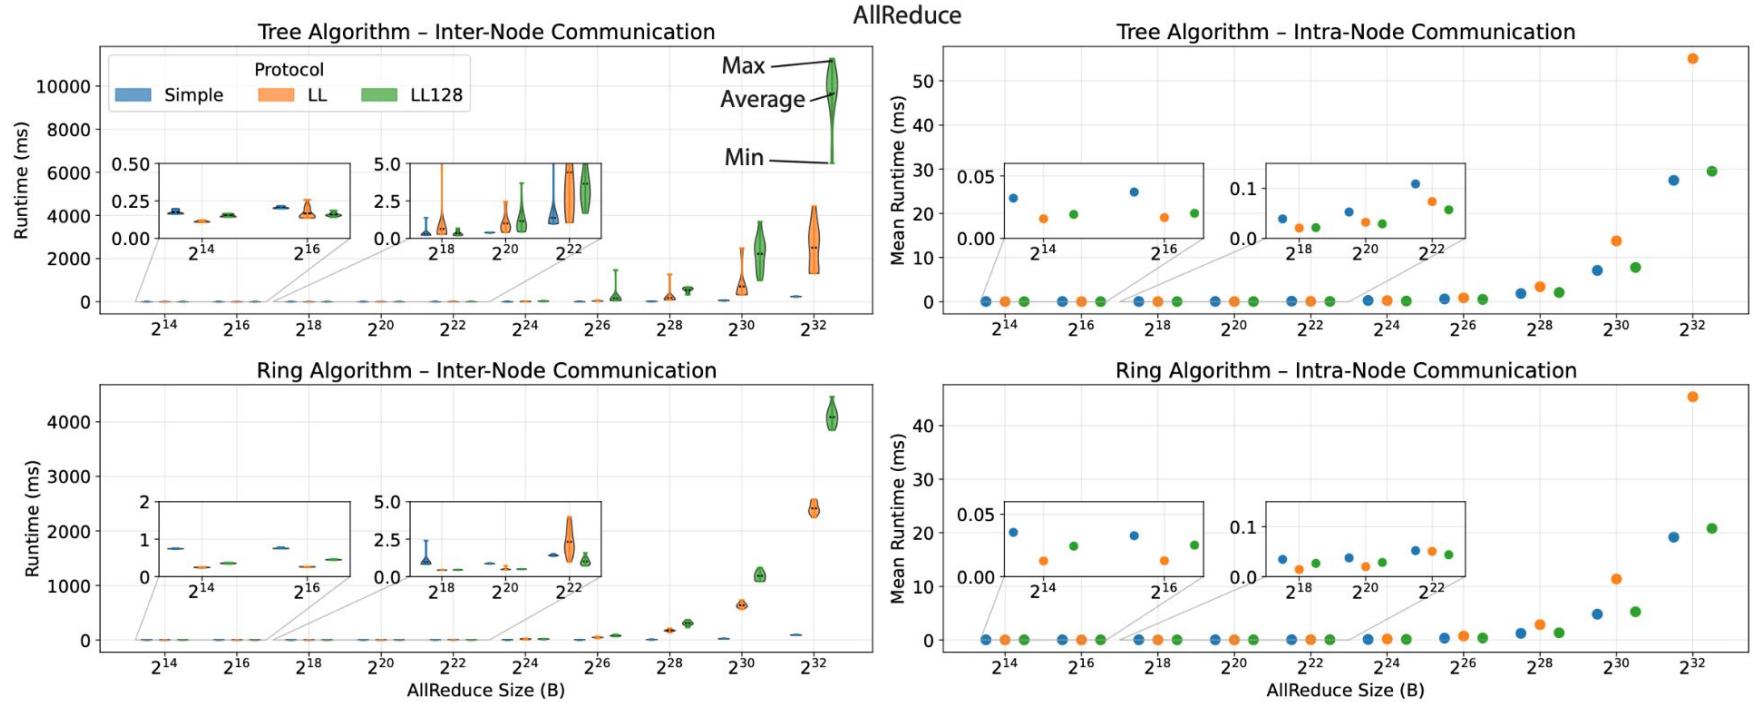


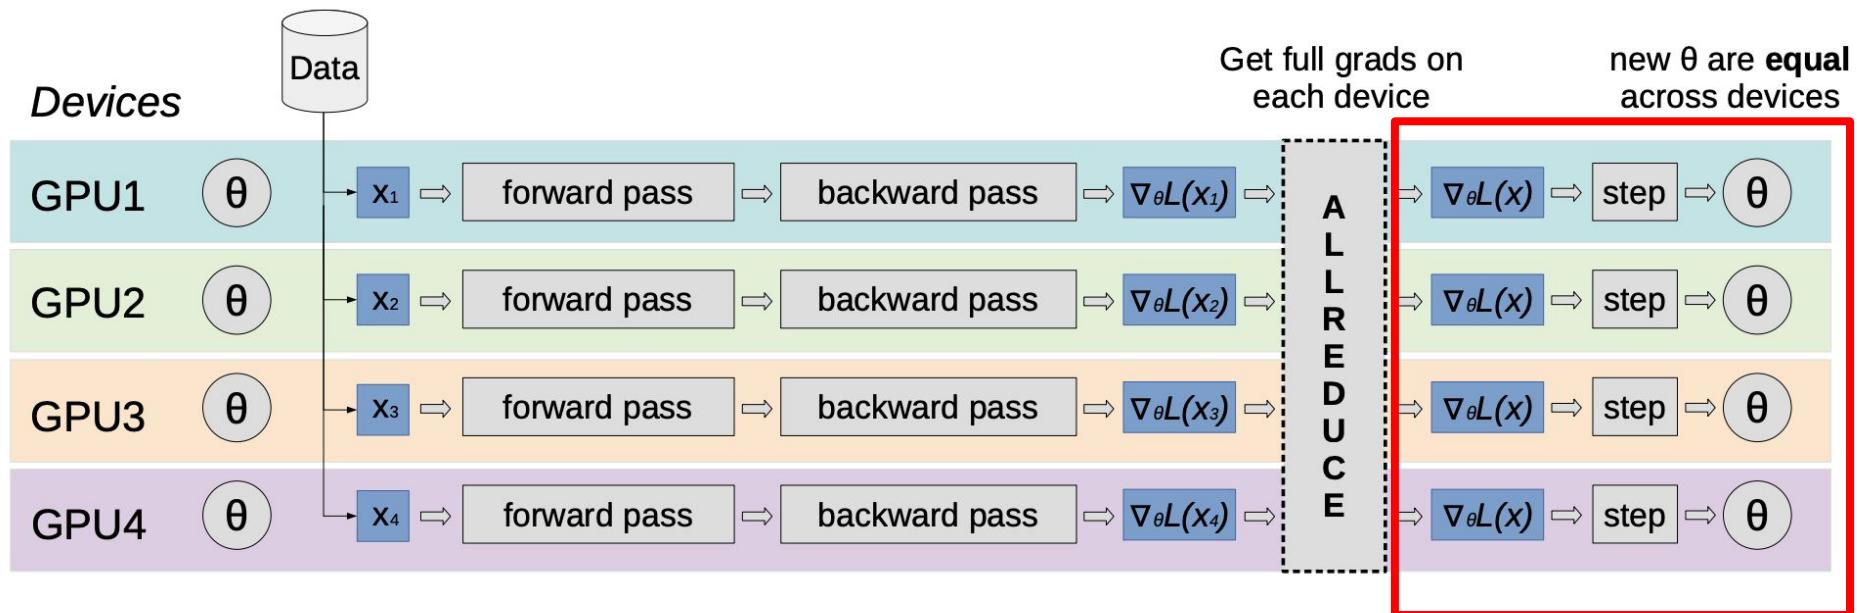
Fig. 6. Runtime comparison of protocols for Ring and Tree AllReduce when running inter- and intra-node. Each data point consists of 20 runs with a warm-up phase. For intra-node communication we report only the median value for readability as the variance is very low.

# AllReduce calculations

```
(main) root@C.27273295:/workspace/nccl-tests$ ./build/all_reduce_perf -g 2
# nccl-tests version 2.17.4 nccl-headers=21701 nccl-library=21701
# Collective test starting: all_reduce_perf
# nThread 1 nGpus 2 minBytes 33554432 maxBytes 33554432 step: 1048576(bytes) warmup iters: 1 iters: 20 agg iters: 1
validation: 1 graph: 0
#
# Using devices
# Rank 0 Group 0 Pid 12457 on bd14101a12c9 device 0 [0000:41:00] Tesla V100-PCIE-16GB
# Rank 1 Group 0 Pid 12457 on bd14101a12c9 device 1 [0000:61:00] Tesla V100-PCIE-16GB
#
#
#          size        count      type    redop     root      time      out-of-place           in-place
#          (B)      (elements)          float      sum      -1      (us)      algbw      busbw #wrong      time      algbw      busbw #wro
ng
#          (B)      (elements)
# 33554432      8388608      float      sum      -1  2943.25      11.40      11.40      0  2953.54      11.36      11.36
0
# Out of bounds values : 0 OK
# Avg bus bandwidth   : 11.3806
#
# Collective test concluded: all_reduce_perf
(main) root@C.27273295:/workspace/nccl-tests$
```

2xV100, PCIe, 11.36 GB/s

# Optimizers



# Optimizers

$$v_{k+1} = \beta_1 v_k + (1 - \beta_1) \nabla f(x_k)$$

$$G_{k+1} = \beta_2 G_k + (1 - \beta_2) (\nabla f(x_k))^2$$

$$x_{k+1} = x_k - \frac{\alpha}{\sqrt{G_{k+1} + \varepsilon}} v_{k+1}.$$

Adam Optimizer

# Optimizers

$$\begin{aligned}v_{k+1} &= \beta_1 v_k + (1 - \beta_1) \nabla f(x_k) \\G_{k+1} &= \beta_2 G_k + (1 - \beta_2) (\nabla f(x_k))^2 \\x_{k+1} &= x_k - \frac{\alpha}{\sqrt{G_{k+1} + \epsilon}} v_{k+1}.\end{aligned}$$

Adam Optimizer

# Optimizers

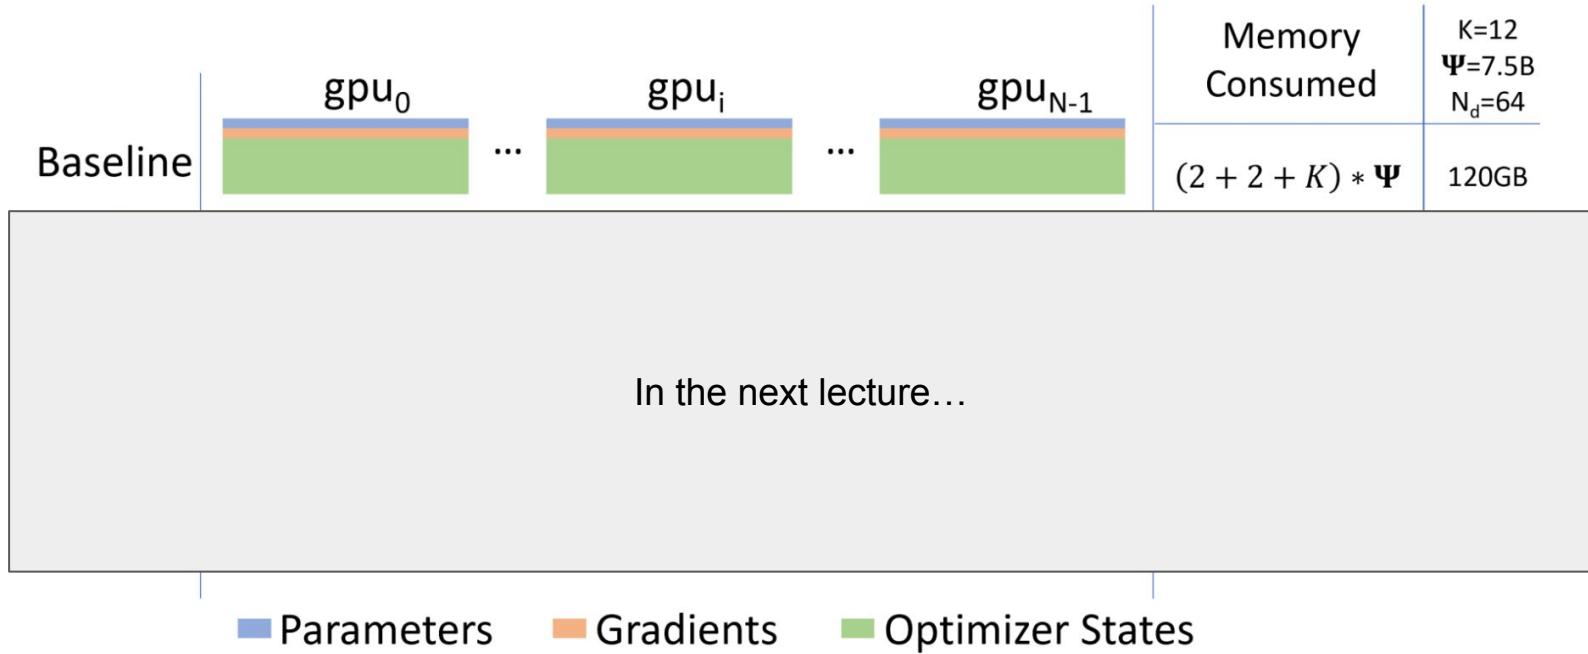
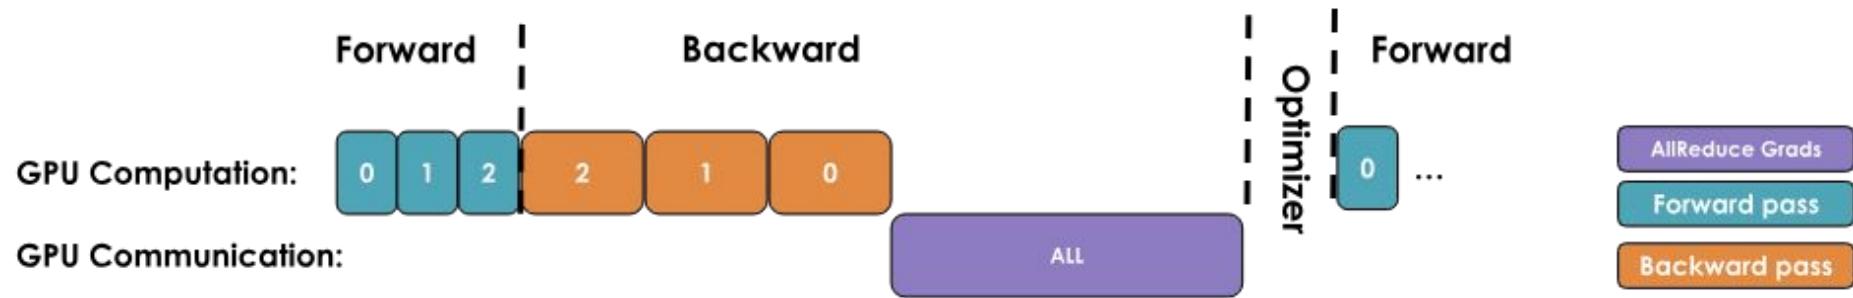
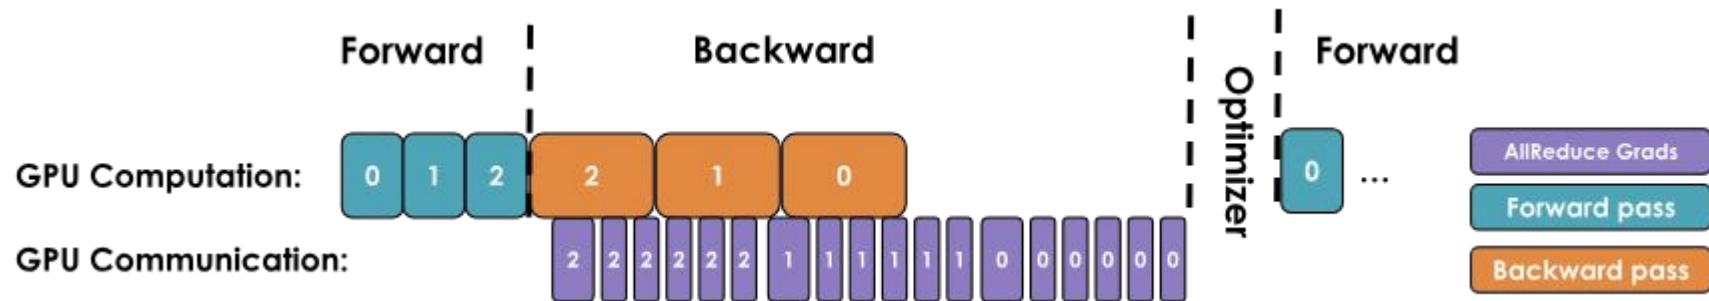


Figure 1: Comparing the per-device memory consumption of model states, with three stages of *ZeRO*-DP optimizations.  $\Psi$  denotes model size (number of parameters),  $K$  denotes the memory multiplier of optimizer states, and  $N_d$  denotes DP degree. In the example, we assume a model size of  $\Psi = 7.5B$  and DP of  $N_d = 64$  with  $K = 12$  based on mixed-precision training with Adam optimizer.

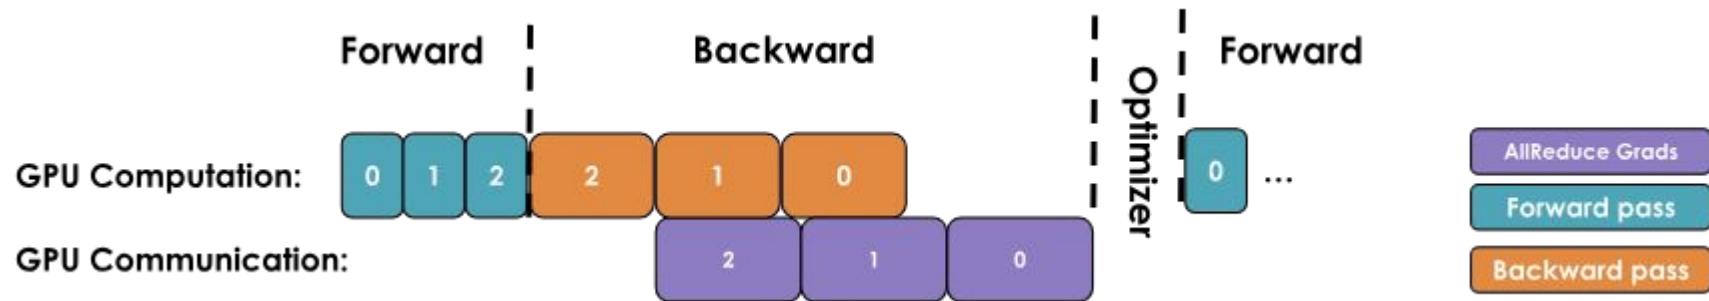
# AllReduce + Optimizer step



# AllReduce + Optimizer step



# AllReduce + Optimizer step



# AllReduce + Optimizer step

```
class DataParallelNaive(nn.Module):
    def __init__(self, module):
        ...
        self.register_backward_hook(self._allreduce_grads)

    def forward(self, *inputs, **kwargs):
        return self.module(*inputs, **kwargs)

    def register_backward_hook(self, hook):
        for p in self.module.parameters():
            if p.requires_grad is True:
                p.register_hook(hook)

    def _allreduce_grads(self, grad):
        dist.all_reduce(grad, op=dist.ReduceOp.SUM, group=pgm.process_group_manager.cp_dp_group)
        grad /= pgm.process_group_manager.cp_dp_world_size
        return grad
```

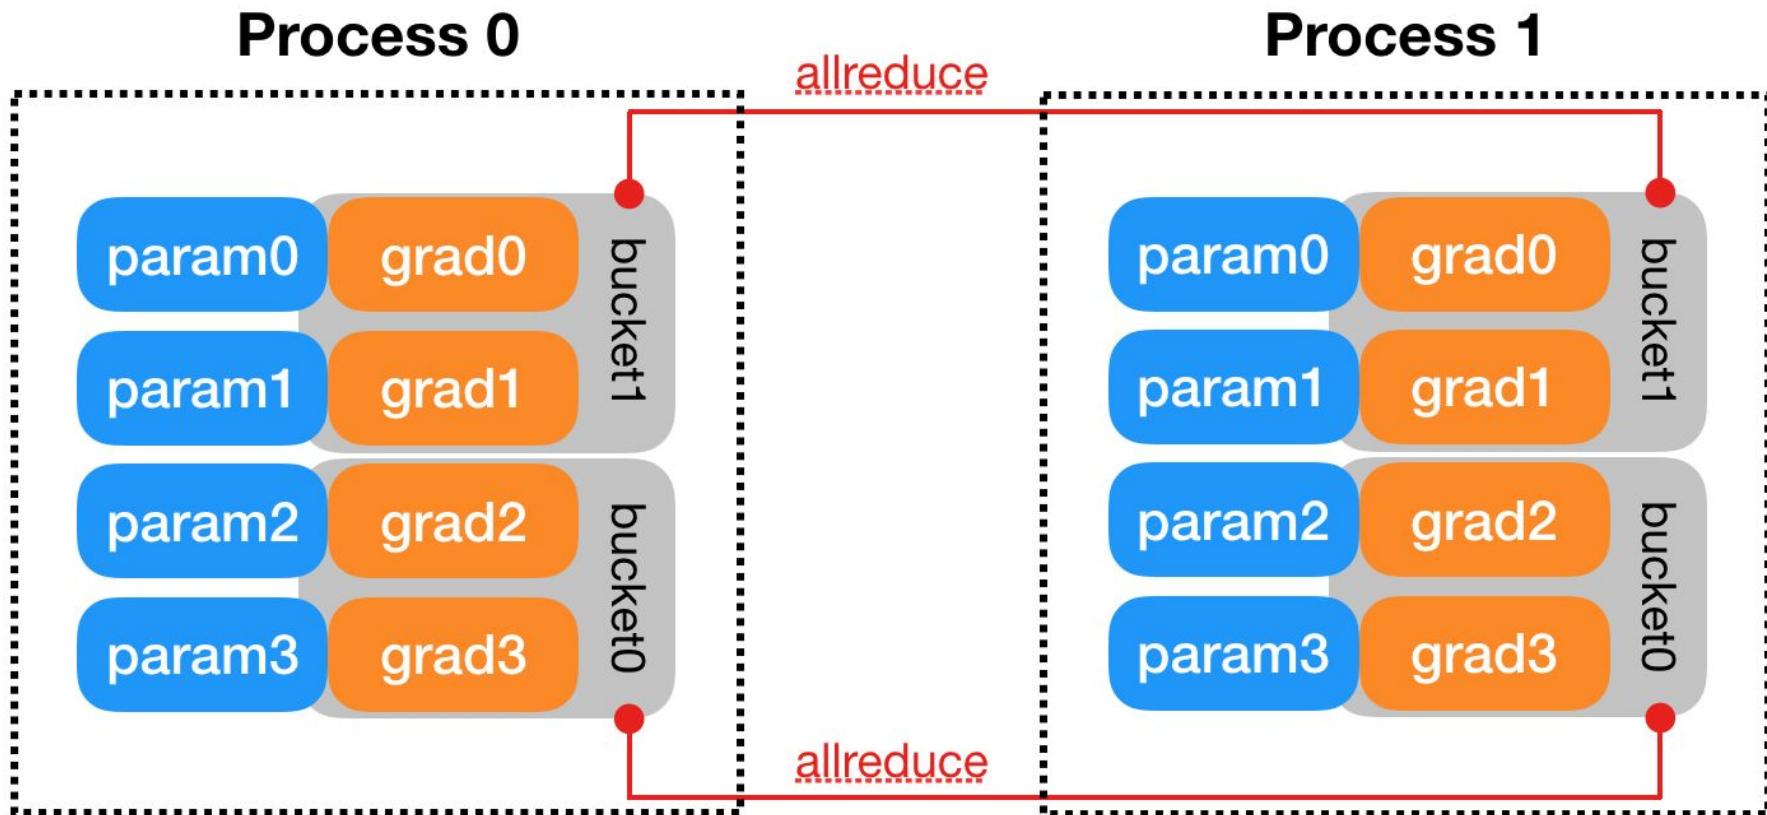
# Pytorch DDP

```
from torch.nn.parallel import DistributedDataParallel as DDP

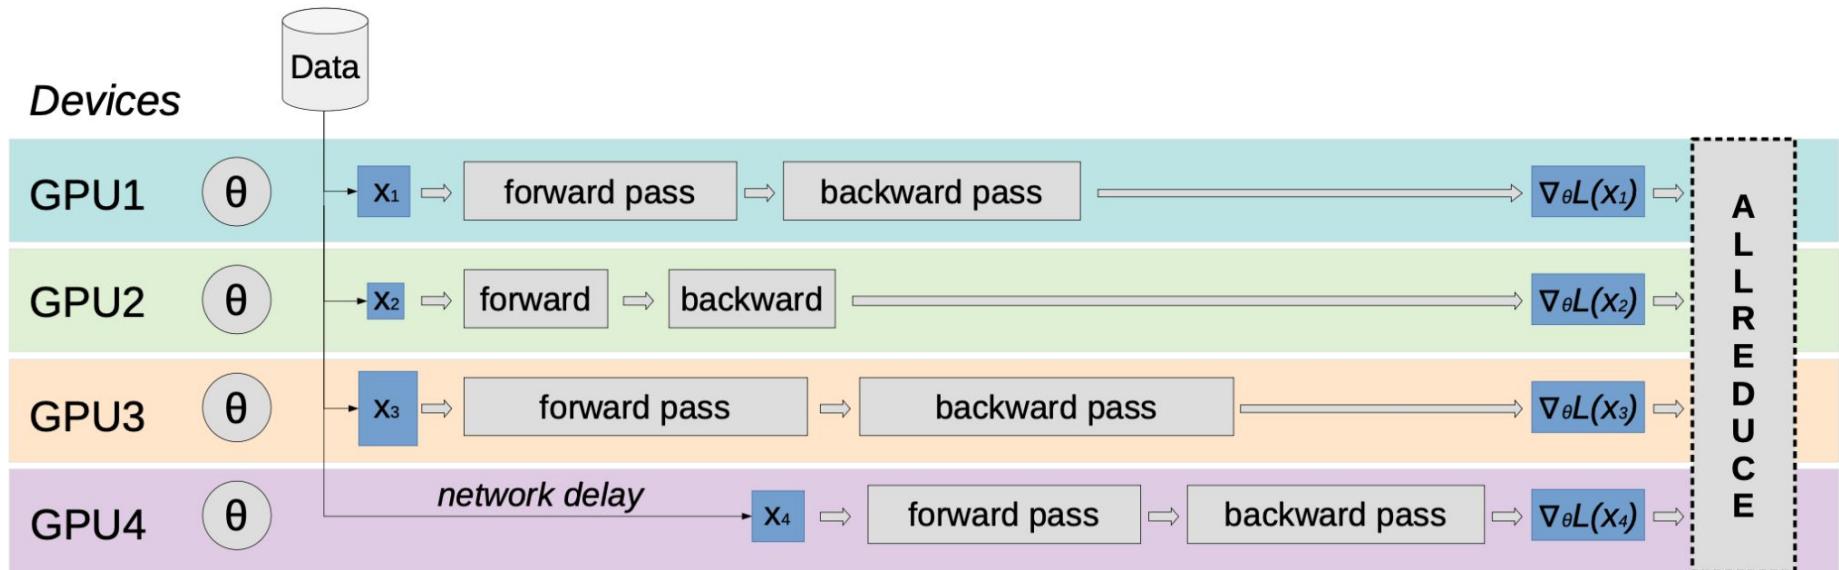
def example(rank, world_size):
    # create default process group
    dist.init_process_group("gloo", rank=rank, world_size=world_size)
    # create local model
    model = nn.Linear(10, 10).to(rank)
    # construct DDP model
    ddp_model = DDP(model, device_ids=[rank])
    # define loss function and optimizer
    loss_fn = nn.MSELoss()
    optimizer = optim.SGD(ddp_model.parameters(), lr=0.001)

    # forward pass
    outputs = ddp_model(torch.randn(20, 10).to(rank))
    labels = torch.randn(20, 10).to(rank)
    # backward pass
    loss_fn(outputs, labels).backward()
    # update parameters
    optimizer.step()
```

# Pytorch DDP

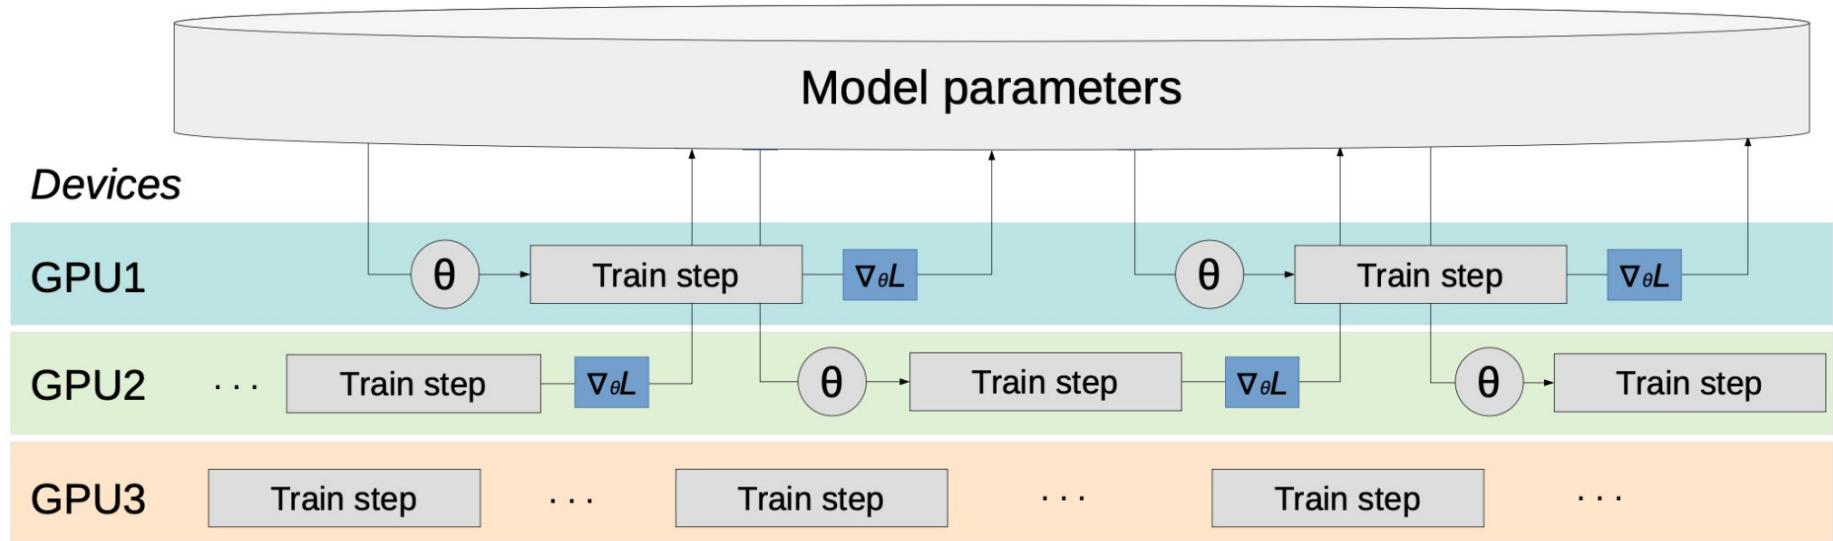


# Heterogeneity issue

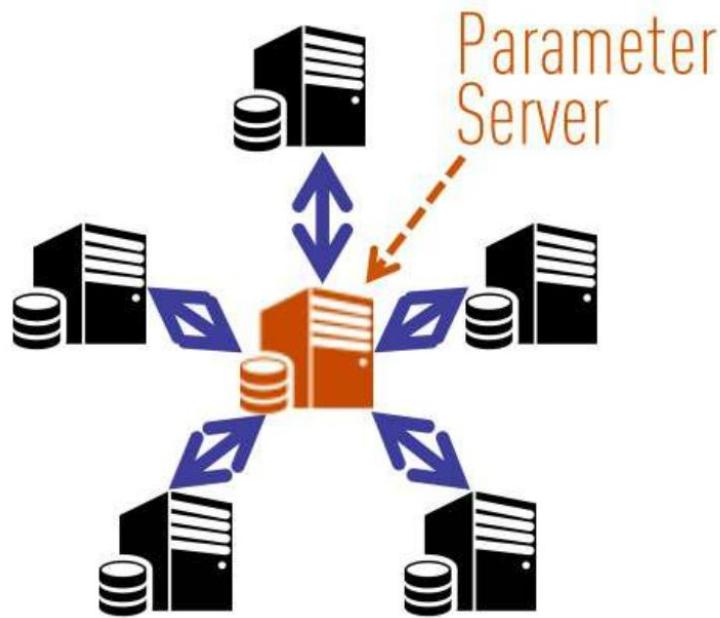


# Heterogeneity issue

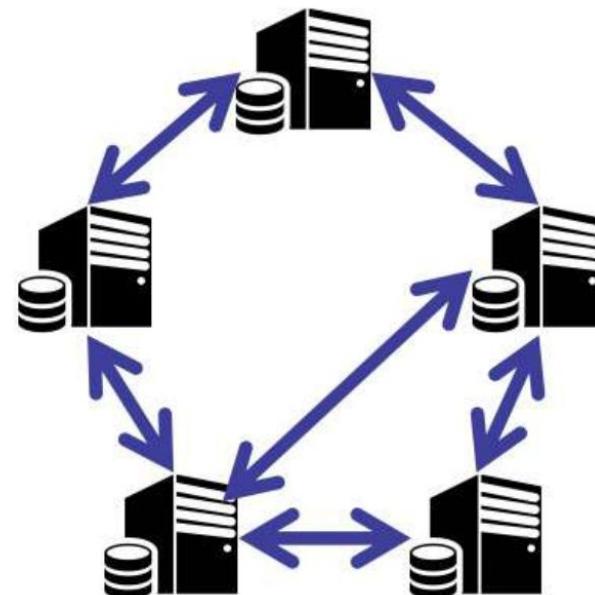
Idea: remove synchronization step altogether, use parameter server



# Decentralized Training with Gossip



**(a) Centralized Topology**



**(b) Decentralized Topology**

# Q&A

# Семинар

# AllReduce calculations

```
(main) root@C.27273295:/workspace/nccl-tests$ ./build/all_reduce_perf -g 2
# nccl-tests version 2.17.4 nccl-headers=21701 nccl-library=21701
# Collective test starting: all_reduce_perf
# nThread 1 nGpus 2 minBytes 33554432 maxBytes 33554432 step: 1048576(bytes) warmup iters: 1 iters: 20 agg iters: 1
validation: 1 graph: 0
#
# Using devices
# Rank 0 Group 0 Pid 12457 on bd14101a12c9 device 0 [0000:41:00] Tesla V100-PCIE-16GB
# Rank 1 Group 0 Pid 12457 on bd14101a12c9 device 1 [0000:61:00] Tesla V100-PCIE-16GB
#
#
#          size        count      type    redop     root      time   out-of-place           in-place
#          (B)      (elements)          float      sum      -1      (us)    algbw    busbw #wrong      time    algbw    busbw #wro
ng
#          (B)      (elements)
# 33554432      8388608      float      sum      -1  2943.25    11.40    11.40      0  2953.54    11.36    11.36
0
# Out of bounds values : 0 OK
# Avg bus bandwidth    : 11.3806
#
# Collective test concluded: all_reduce_perf
(main) root@C.27273295:/workspace/nccl-tests$
```

2xV100, PCIe, 11.36 GB/s